

DATA ANALYSIS OF DIABETES DATABASE ***

```
#import libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings("ignore")#call the class for adaboost
from sklearn.ensemble import AdaBoostClassifier
```

```
df=pd.read_csv("diabetes.csv")
```

```
df.head()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigre
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	

```
#check null values
df.isnull().sum()
```

Pregnancies	0
Glucose	0
BloodPressure	0
SkinThickness	0
Insulin	0
BMI	0
DiabetesPedigreeFunction	0
Age	0
Outcome	0
dtype: int64	

```
#check inormation
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column              Non-Null Count  Dtype
---  -
0   Pregnancies         768 non-null    int64
1   Glucose              768 non-null    int64
```

```

2   BloodPressure      768 non-null    int64
3   SkinThickness      768 non-null    int64
4   Insulin            768 non-null    int64
5   BMI                768 non-null    float64
6   DiabetesPedigreeFunction 768 non-null    float64
7   Age                768 non-null    int64
8   Outcome            768 non-null    int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB

```

```
#display all rows
```

```
pd.set_option('display.max_rows',None)
```

```
#display all columns
```

```
pd.set_option('display.max_columns',None)
```

```
#seperate the independent (input X) and dependent (output Y) variable
```

```
X=df.drop("Outcome",axis=1)
```

```
Y=df["Outcome"]
```

```
#call class for training and testing
```

```
from sklearn.model_selection import train_test_split
```

```
#split data into training and testing
```

```
X_train,X_test,Y_train,Y_test=train_test_split(X,Y,test_size=0.3, random_state=1)
```

```
#create a function
```

```
def create_model(model):
    model.fit(X_train,Y_train)#to train
    y_pred=model.predict(X_test)
    print(classification_report(Y_test,y_pred))
    return model
```

```
#as we have used classification report we hhave to call the class for it
```

```
from sklearn.metrics import classification_report
```

```
#call class for logistic regression
```

```
from sklearn.linear_model import LogisticRegression
```

```
#create object for LogisticRegression class
```

```
lr=LogisticRegression()
```

```
#call the function "create_model"(for training and testing)
```

```
create_model(lr)
```

```

              precision    recall  f1-score   support

0               0.79         0.90         0.84         146
1               0.78         0.58         0.66          85

```

accuracy			0.78	231
macro avg	0.78	0.74	0.75	231
weighted avg	0.78	0.78	0.78	231

```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
intercept_scaling=1, l1_ratio=None, max_iter=100,
multi_class='auto', n_jobs=None, penalty='l2',
random_state=None, solver='lbfgs', tol=0.0001, verbose=0,
warm_start=False)
```

```
LogisticRegression()
```

```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
intercept_scaling=1, l1_ratio=None, max_iter=100,
multi_class='auto', n_jobs=None, penalty='l2',
random_state=None, solver='lbfgs', tol=0.0001, verbose=0,
warm_start=False)
```

```
#call the class for DecisionTreeClassifier
from sklearn.tree import DecisionTreeClassifier
```

```
#create object of DecisionTreeClassifier class
dt=DecisionTreeClassifier()
```

```
#call function "create_model"(for training and testing)
dt=create_model(dt)
```

	precision	recall	f1-score	support
0	0.73	0.82	0.77	146
1	0.60	0.47	0.53	85

accuracy			0.69	231
macro avg	0.66	0.64	0.65	231
weighted avg	0.68	0.69	0.68	231

```
#to check the information gain (if information gain is high then the column will be :
dt.feature_importances_
```

```
array([0.0707469 , 0.26611293, 0.16709624, 0.0400096 , 0.05283241,
0.18573917, 0.10574786, 0.11171488])
```

```
#call the class for tree
from sklearn import tree
```


```
#keep all the input variables in features
features=X.columns
```

```
#draw the tree
```

```
fig=plt.figure(figsize=(20,20))
```

```
fig=plt.figure(figsize=(25,30))
```

```
_tree.plot_tree(dt,feature_names=features,filled=True)
```



```
#(1) max_depth
```

```
#we create a different object of DecisionTreeClassifier and pass parameter for max_depth
```

```
dt1=DecisionTreeClassifier(max_depth=5) #not more than 8
```

```
#call function "create_model"(for training and testing)
```

```
dt1=create_model(dt1)
```

```
precision    recall  f1-score   support
```

0	0.80	0.82	0.81	146
1	0.68	0.65	0.66	85
accuracy			0.76	231
macro avg	0.74	0.73	0.74	231
weighted avg	0.76	0.76	0.76	231

#it is giving good recall when the max_depth is 5 i.e 0.67

#(2) min_sample_leaf

#we create a different object of DecisionTreeClassifier and pass parameter for min_s:
dt2=DecisionTreeClassifier(min_samples_leaf=45) #min 50

#call function "create_model"(for training and testing)
dt2=create_model(dt2)

	precision	recall	f1-score	support
0	0.83	0.86	0.84	146
1	0.74	0.71	0.72	85
accuracy			0.80	231
macro avg	0.79	0.78	0.78	231
weighted avg	0.80	0.80	0.80	231

#(3) entropy

#we create a different object of DecisionTreeClassifier and pass parameter for min_s:
dt3=DecisionTreeClassifier(min_samples_leaf=40,criterion="entropy")

#call function "create_model"(for training and testing)
dt3=create_model(dt3)

	precision	recall	f1-score	support
0	0.82	0.85	0.83	146
1	0.72	0.67	0.70	85
accuracy			0.78	231
macro avg	0.77	0.76	0.76	231
weighted avg	0.78	0.78	0.78	231

it is not giving nice recall that is the model is getting overfit so the max-depth is the best technique for training as it is giving good recall but recall is not still that good so we will apply next algorithm

Double-click (or enter) to edit

SUPPORT VECTOR MACHINE (SVM)

```
#call class form svm
from sklearn.svm import LinearSVC

#create object for LinearSVC
svc=LinearSVC(random_state=1,C=0.0012) #c is paramete for adding error

#call function "create_model"(for training and testing)
create_model(svc)
```

	precision	recall	f1-score	support
0	0.68	0.88	0.77	146
1	0.59	0.31	0.40	85
accuracy			0.67	231
macro avg	0.64	0.59	0.59	231
weighted avg	0.65	0.67	0.63	231

```
LinearSVC(C=0.0012, class_weight=None, dual=True, fit_intercept=True,
          intercept_scaling=1, loss='squared_hinge', max_iter=1000,
          multi_class='ovr', penalty='l2', random_state=1, tol=0.0001,
          verbose=0)
```

```
LinearSVC(C=0.0012, random_state=1)
```

```
LinearSVC(C=0.0012, class_weight=None, dual=True, fit_intercept=True,
          intercept_scaling=1, loss='squared_hinge', max_iter=1000,
          multi_class='ovr', penalty='l2', random_state=1, tol=0.0001,
          verbose=0)
```

ENSEMBLING TECHNIQUE

1. ADABOOST

```
#call the class for adaboost
from sklearn.ensemble import AdaBoostClassifier
```

```
#create object of Adaboost
ada=AdaBoostClassifier(n_estimators=100)
```

```
#call function "create_model"
create_model(ada)
```

	precision	recall	f1-score	support
0	0.82	0.88	0.85	146
1	0.77	0.66	0.71	85

accuracy			0.80	231
macro avg	0.79	0.77	0.78	231
weighted avg	0.80	0.80	0.80	231

```
AdaBoostClassifier(algorithm='SAMME.R', base_estimator=None, learning_rate=1.0,
                    n_estimators=100, random_state=None)
```

```
AdaBoostClassifier(n_estimators=100)
```

```
AdaBoostClassifier(algorithm='SAMME.R', base_estimator=None, learning_rate=1.0,
                    n_estimators=100, random_state=None)
```

2. GRADIENT BOOST

```
#call the class for GradientBoost
from sklearn.ensemble import GradientBoostingClassifier
```

```
#create object of GradientBoostingClassifier
gbc=GradientBoostingClassifier(n_estimators=100)
```

```
#call function"create_model"
create_model(gbc)
```

	precision	recall	f1-score	support
0	0.81	0.88	0.85	146
1	0.76	0.65	0.70	85

accuracy			0.80	231
macro avg	0.79	0.77	0.77	231
weighted avg	0.79	0.80	0.79	231

```
GradientBoostingClassifier(ccp_alpha=0.0, criterion='friedman_mse', init=None,
                           learning_rate=0.1, loss='deviance', max_depth=3,
                           max_features=None, max_leaf_nodes=None,
                           min_impurity_decrease=0.0, min_impurity_split=None,
                           min_samples_leaf=1, min_samples_split=2,
                           min_weight_fraction_leaf=0.0, n_estimators=100,
                           n_iter_no_change=None, presort='deprecated',
                           random_state=None, subsample=1.0, tol=0.0001,
                           validation_fraction=0.1, verbose=0,
                           warm_start=False)
```

```
GradientBoostingClassifier()
```

```
GradientBoostingClassifier(ccp_alpha=0.0, criterion='friedman_mse', init=None,
                           learning_rate=0.1, loss='deviance', max_depth=3,
                           max_features=None, max_leaf_nodes=None,
                           min_impurity_decrease=0.0, min_impurity_split=None,
                           min_samples_leaf=1, min_samples_split=2,
                           min_weight_fraction_leaf=0.0, n_estimators=100,
                           n_iter_no_change=None, presort='deprecated',
                           random_state=None, subsample=1.0, tol=0.0001,
```

```
validation_fraction=0.1, verbose=0,
warm_start=False)
```

3. NAIVE

```
#create model list
model_list=[("Logistic",lr),("DecisonTree",dt),("DecisionTreeEntropy",dt3)]

#call class for hard voting
from sklearn.ensemble import VotingClassifier

# (A) HARD VOTING

f VotingClassifier
ier(estimators=model_list) #by default it takes hard voting so if we are not specifyi

#call function "create_model"
create_model(vc)
```

	precision	recall	f1-score	support
0	0.80	0.88	0.84	146
1	0.75	0.61	0.68	85
accuracy			0.78	231
macro avg	0.77	0.75	0.76	231
weighted avg	0.78	0.78	0.78	231

```
VotingClassifier(estimators=[('Logistic',
                             LogisticRegression(C=1.0, class_weight=None,
                                                  dual=False, fit_intercept=True,
                                                  intercept_scaling=1,
                                                  l1_ratio=None, max_iter=100,
                                                  multi_class='auto',
                                                  n_jobs=None, penalty='l2',
                                                  random_state=None,
                                                  solver='lbfgs', tol=0.0001,
                                                  verbose=0, warm_start=False)),
                             ('DecisonTree',
                              DecisionTreeClassifier(ccp_alpha=0.0,
                                                       class_weight=None,
                                                       cr...
                              DecisionTreeClassifier(ccp_alpha=0.0,
                                                       class_weight=None,
                                                       criterion='entropy',
                                                       max_depth=None,
                                                       max_features=None,
                                                       max_leaf_nodes=None,
                                                       min_impurity_decrease=0.0,
                                                       min_impurity_split=None,
                                                       min_samples_leaf=40,
                                                       min_samples_split=2,
                                                       min_weight_fraction_leaf=0.0,
```



```

presort='deprecated',
random_state=None,
splitter='best'))],
flatten_transform=True, n_jobs=None, voting='hard',
weights=None)

```

#4. BOOTSTRAPING

```

#create model list
model_list=[("Logistic",lr),("DecisonTree",dt),("DecisionTreeEntropy",dt3)]

#call class for BaggingClassifier
from sklearn.ensemble import BaggingClassifier

object of BaggingClassifier
BaggingClassifier(LogisticRegression(),n_estimators=10,max_samples=10,random_state=1) #by

```

(B) PASTING

```

#create object of BaggingClassifier
bc1=BaggingClassifier(LogisticRegression(),n_estimators=10,max_samples=10,random_state=1)

```

```

#call function "create_model"
create_model(bc1)

```

	precision	recall	f1-score	support
0	0.78	0.88	0.83	146
1	0.74	0.56	0.64	85
accuracy			0.77	231
macro avg	0.76	0.72	0.73	231
weighted avg	0.76	0.77	0.76	231

```

BaggingClassifier(base_estimator=LogisticRegression(C=1.0, class_weight=None,
dual=False,
fit_intercept=True,
intercept_scaling=1,
l1_ratio=None, max_iter=100,
multi_class='auto',
n_jobs=None, penalty='l2',
random_state=None,
solver='lbfgs', tol=0.0001,
verbose=0,
warm_start=False),
bootstrap=False, bootstrap_features=False, max_features=1.0,
max_samples=10, n_estimators=10, n_jobs=None, oob_score=False,
random_state=1, verbose=0, warm_start=False)

```

(C) RANDOM FOREST

```
#call class for RandomForestClassifier
from sklearn.ensemble import RandomForestClassifier

#create object of RandomForestClassifier
rf=RandomForestClassifier(n_estimators=10,max_features=7,random_state=1)
```

```
#call function "create_model"
create_model(rf)
```

	precision	recall	f1-score	support
0	0.76	0.90	0.82	146
1	0.75	0.52	0.61	85
accuracy			0.76	231
macro avg	0.75	0.71	0.72	231
weighted avg	0.76	0.76	0.75	231

```
RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
                        criterion='gini', max_depth=None, max_features=7,
                        max_leaf_nodes=None, max_samples=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=1, min_samples_split=2,
                        min_weight_fraction_leaf=0.0, n_estimators=10,
                        n_jobs=None, oob_score=False, random_state=1, verbose=0,
                        warm_start=False)
```

FEATURE SELECTION TECHNIQUE

A . ANOVA TEST

```
from sklearn.feature_selection import f_regression #f_regression is for anova test
from sklearn.feature_selection import SelectKBest
#selectKBest is important object
```

```
#create object for SlectKBest class
anova=SelectKBest(score_func=f_regression,k=6) #k is no of columns/feature which retr
#anov auser define object
```

```
#train data and return 10 best column and store in X_train_imp
X_train_imp=anova.fit_transform(X_train,Y_train)
```

```
#test the data
X_test_imp=anova.transform(X_test)
```

```
#check which features are selected and which are rejected
anova.get_support() #return ans in boolean type 0 means false 1 means true
```

```
array([ True,  True, False, False,  True,  True,  True,  True])
```

```
#again create object of logistic regression after anova test
```

```
lr1=LogisticRegression()
```

```
#train model
```

```
lr1.fit(X_train_imp,Y_train)
```

```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                    intercept_scaling=1, l1_ratio=None, max_iter=100,
                    multi_class='auto', n_jobs=None, penalty='l2',
                    random_state=None, solver='lbfgs', tol=0.0001, verbose=0,
                    warm_start=False)
```

```
#check the score
```

```
lr1.score(X_test_imp,Y_test)
```

```
0.7748917748917749
```

```
#CHI SQUARE TEST
```

```
# before applying chi square test to check which if column has non negative
for col in X:
```

```
    print("column name:",col)
```

```
    print(df[col].unique())
```

```
40.6 47.9 50. 25.2 40.9 37.2 44.2 29.9 31.9 28.4 43.5 32.7 67.1 45.
34.9 27.7 35.9 22.6 33.1 30.4 52.3 24.3 22.9 34.8 30.9 40.1 23.9 37.5
35.5 42.8 42.6 41.8 35.8 37.8 28.8 23.6 35.7 36.7 45.2 44. 46.2 35.
43.6 44.1 18.4 29.2 25.9 32.1 36.3 40. 25.1 27.5 45.6 27.8 24.9 25.3
37.9 27. 26. 38.7 20.8 36.1 30.7 32.3 52.9 21. 39.7 25.5 26.2 19.3
38.1 23.5 45.5 23.1 39.9 36.8 21.8 41. 42.2 34.4 27.2 36.5 29.8 39.2
38.4 36.2 48.3 20. 22.3 45.7 23.7 22.1 42.1 42.4 18.2 26.4 45.3 37.
24.5 32.2 59.4 21.2 26.7 30.2 46.1 41.3 38.8 35.2 42.3 40.7 46.5 33.5
37.3 30.3 26.3 21.7 36.4 28.5 26.9 38.6 31.3 19.5 20.1 40.8 23.4 28.3
38.9 57.3 35.6 49.6 44.6 24.1 44.5 41.2 49.3 46.3]
```

```
column name: DiabetesPedigreeFunction
```

```
[0.627 0.351 0.672 0.167 2.288 0.201 0.248 0.134 0.158 0.232 0.191 0.537
1.441 0.398 0.587 0.484 0.551 0.254 0.183 0.529 0.704 0.388 0.451 0.263
0.205 0.257 0.487 0.245 0.337 0.546 0.851 0.267 0.188 0.512 0.966 0.42
0.665 0.503 1.39 0.271 0.696 0.235 0.721 0.294 1.893 0.564 0.586 0.344
0.305 0.491 0.526 0.342 0.467 0.718 0.962 1.781 0.173 0.304 0.27 0.699
0.258 0.203 0.855 0.845 0.334 0.189 0.867 0.411 0.583 0.231 0.396 0.14
0.391 0.37 0.307 0.102 0.767 0.237 0.227 0.698 0.178 0.324 0.153 0.165
0.443 0.261 0.277 0.761 0.255 0.13 0.323 0.356 0.325 1.222 0.179 0.262
0.283 0.93 0.801 0.207 0.287 0.336 0.247 0.199 0.543 0.192 0.588 0.539
0.22 0.654 0.223 0.759 0.26 0.404 0.186 0.278 0.496 0.452 0.403 0.741
0.361 1.114 0.457 0.647 0.088 0.597 0.532 0.703 0.159 0.268 0.286 0.318
0.272 0.572 0.096 1.4 0.218 0.085 0.399 0.432 1.189 0.687 0.137 0.637
0.833 0.229 0.817 0.204 0.368 0.743 0.722 0.256 0.709 0.471 0.495 0.18
0.542 0.773 0.678 0.719 0.382 0.319 0.19 0.956 0.084 0.725 0.299 0.244
0.745 0.615 1.321 0.64 0.142 0.374 0.383 0.578 0.136 0.395 0.187 0.905
0.15 0.874 0.236 0.787 0.407 0.605 0.151 0.289 0.355 0.29 0.375 0.164
0.431 0.742 0.514 0.464 1.224 1.072 0.805 0.209 0.666 0.101 0.198 0.652
2.329 0.089 0.645 0.238 0.394 0.293 0.479 0.686 0.831 0.582 0.446 0.402
1.318 0.329 1.213 0.427 0.282 0.143 0.38 0.284 0.249 0.926 0.557 0.092
0.655 1.353 0.612 0.2 0.226 0.997 0.933 1.101 0.078 0.24 1.136 0.128
0.422 0.251 0.677 0.296 0.454 0.744 0.881 0.28 0.259 0.619 0.808 0.34
0.424 0.757 0.612 0.602 0.52 0.412 0.84 0.820 0.156 0.315 0.326 1.301]
```

```

0.454 0.757 0.613 0.692 0.52 0.412 0.84 0.859 0.156 0.215 0.326 1.391
0.875 0.313 0.433 0.626 1.127 0.315 0.345 0.129 0.527 0.197 0.731 0.148
0.123 0.127 0.122 1.476 0.166 0.932 0.343 0.893 0.331 0.472 0.673 0.389
0.485 0.349 0.279 0.346 0.252 0.243 0.58 0.559 0.302 0.569 0.378 0.385
0.499 0.306 0.234 2.137 1.731 0.545 0.225 0.816 0.528 0.509 1.021 0.821
0.947 1.268 0.221 0.66 0.239 0.949 0.444 0.463 0.803 1.6 0.944 0.196
0.241 0.161 0.135 0.376 1.191 0.702 0.674 1.076 0.534 1.095 0.554 0.624
0.219 0.507 0.561 0.421 0.516 0.264 0.328 0.233 0.108 1.138 0.147 0.727
0.435 0.497 0.23 0.955 2.42 0.658 0.33 0.51 0.285 0.415 0.381 0.832
0.498 0.212 0.364 1.001 0.46 0.733 0.416 0.705 1.022 0.269 0.6 0.571
0.607 0.17 0.21 0.126 0.711 0.466 0.162 0.419 0.63 0.365 0.536 1.159
0.629 0.292 0.145 1.144 0.174 0.547 0.163 0.738 0.314 0.968 0.409 0.297
0.525 0.154 0.771 0.107 0.493 0.717 0.917 0.501 1.251 0.735 0.804 0.661
0.549 0.825 0.423 1.034 0.16 0.341 0.68 0.591 0.3 0.121 0.502 0.401
0.601 0.748 0.338 0.43 0.892 0.813 0.693 0.575 0.371 0.206 0.417 1.154
0.925 0.175 1.699 0.682 0.194 0.4 0.1 1.258 0.482 0.138 0.593 0.878
0.157 1.282 0.141 0.246 1.698 1.461 0.347 0.362 0.393 0.144 0.732 0.115
0.465 0.649 0.871 0.149 0.695 0.303 0.61 0.73 0.447 0.455 0.133 0.155
1.162 1.292 0.182 1.394 0.217 0.631 0.88 0.614 0.332 0.366 0.181 0.828
0.335 0.856 0.886 0.439 0.253 0.598 0.904 0.483 0.565 0.118 0.177 0.176

0.295 0.441 0.352 0.826 0.97 0.595 0.317 0.265 0.646 0.426 0.56 0.515
0.453 0.785 0.734 1.174 0.488 0.358 1.096 0.408 1.182 0.222 1.057 0.766
0.171]
column name: Age
[50 31 32 21 33 30 26 29 53 54 34 57 59 51 27 41 43 22 38 60 28 45 35 46
 56 37 48 40 25 24 58 42 44 39 36 23 61 69 62 55 65 47 52 66 49 63 67 72
 81 64 70 68]

```

```

#call class for chi square test
from sklearn.feature_selection import chi2

#create object
chi=SelectKBest(score_func=chi2,k=6)

#train data and return 10 best column and store in X_train_imp
X_train_imp=chi.fit_transform(X_train,Y_train)

#check which features are selected and which are rejected
anova.get_support() #return ans in boolean type 0 means false 1 means true

array([ True,  True, False, False,  True,  True,  True,  True])

#again create object of logistic regression after anova test
lr1=LogisticRegression()

#train model
lr1.fit(X_train_imp,Y_train)

LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                    intercept_scaling=1, l1_ratio=None, max_iter=100,
                    multi_class='auto', n_jobs=None, penalty='l2',
                    random_state=None, solver='lbfgs', tol=0.0001, verbose=0,
                    warm_start=False)

```

```
#check the score
lr1.score(X_test_imp,Y_test)
```

```
0.6320346320346321
```

#2. WRAPPER METHOD

```
#declare empty list for the columns
columns=[]
```

```
for col in X:
    columns.append(col)
    X_new=df[columns] #input variable
    X_train,X_test,Y_train,Y_test = train_test_split(X_new,Y,test_size=0.3,random_st:
    #create a object of LogisticRegression
    lin=LogisticRegression()
    #we train the model
    lin.fit(X_train,Y_train)
    #find the score
    score1=lin.score(X_test,Y_test)
    print("Column : ",col, " Score : ",score1)
```


```
Column : Pregnancies Score : 0.6536796536796536
Column : Glucose Score : 0.7532467532467533
Column : BloodPressure Score : 0.7619047619047619
Column : SkinThickness Score : 0.7705627705627706
Column : Insulin Score : 0.7748917748917749
Column : BMI Score : 0.7835497835497836
Column : DiabetesPedigreeFunction Score : 0.7878787878787878
Column : Age Score : 0.7835497835497836
```

SAMPLING TECHNIQUE

```
#to count how many yes and how many no
df["Outcome"].value_counts()
```

```
0    500
1    268
Name: Outcome, dtype: int64
```

```
#display using countplot
sns.countplot(data=df,x="Outcome")
plt.show()
```



There are two types of sampling techniques suppose - yes = 10 (minority class) no = 1000 (majority class)

- 1) Over Sampling :when we add duplicate rows (increasing minority class) it is called over sampling
- 2) Under Sampling :when we delete the rows (reducing majority class) it is called as under sampling

```
#call class for over sampling and under sampling
from imblearn.under_sampling import RandomUnderSampler
from imblearn.over_sampling import RandomOverSampler
```

(1) Over Sampling

```
#count yes and no before applying sampling
pd.Series(Y_train).value_counts()
```

```
0    354
1    183
Name: Outcome, dtype: int64
```

```
#create object of RandomOverSampler
ros= RandomOverSampler()
```

```
#train the data using sampling
X_sample2,y_sample2 = ros.fit_sample(X_train,Y_train)
```

```
#count yes and no after applying sampling
pd.Series(y_sample2).value_counts()
```

```
1    354
0    354
dtype: int64
```

```
#now apply the Decision tree (of classification algorithm) using pruning technique
dt4= DecisionTreeClassifier(max_depth=4)
```

```
#now train the model
dt4.fit(X_sample2,y_sample2)
```

```
DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None, criterion='gini',
                        max_depth=4, max_features=None, max_leaf_nodes=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
```

```
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, presort='deprecated',
random_state=None, splitter='best')
```

```
#now test the model
y_pred=dt4.predict(X_test)
```

```
#classification report
print(classification_report(Y_test,y_pred))
```

	precision	recall	f1-score	support
0	0.82	0.80	0.81	146
1	0.67	0.69	0.68	85
accuracy			0.76	231
macro avg	0.74	0.75	0.75	231
weighted avg	0.76	0.76	0.76	231

▼ (2) Under sampling

```
#create object of RandomUnderSampler
rus=RandomUnderSampler()
```

```
#train the data using sampling
X_sample1,y_sample1 = rus.fit_sample(X_train,Y_train)
```

```
#count yes and no after appyling sampling
pd.Series(y_sample1).value_counts()
```

```
1    183
0    183
dtype: int64
```

```
#now apply the Decision tree (of classification algorithm) using purning technique
dt5= DecisionTreeClassifier(max_depth=3)
```


```
#now train the model
dt5.fit(X_sample1,y_sample1)
```

```
DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None, criterion='gini',
max_depth=3, max_features=None, max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, presort='deprecated',
random_state=None, splitter='best')
```

```
#now test the model
y_pred=dt5.predict(X_test)
```

```
y_pred=cls.predict(X_test)
```

```
#classification report
print(classification_report(Y_test,y_pred))
```



	precision	recall	f1-score	support
0	0.92	0.60	0.72	146
1	0.57	0.91	0.70	85
accuracy			0.71	231
macro avg	0.74	0.75	0.71	231
weighted avg	0.79	0.71	0.71	231